

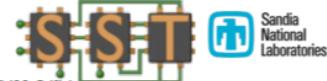
Overview

Created By	 Magic Griffin
Last Edited	@Jul 05, 2019 11:24 PM
Tags	

Brief

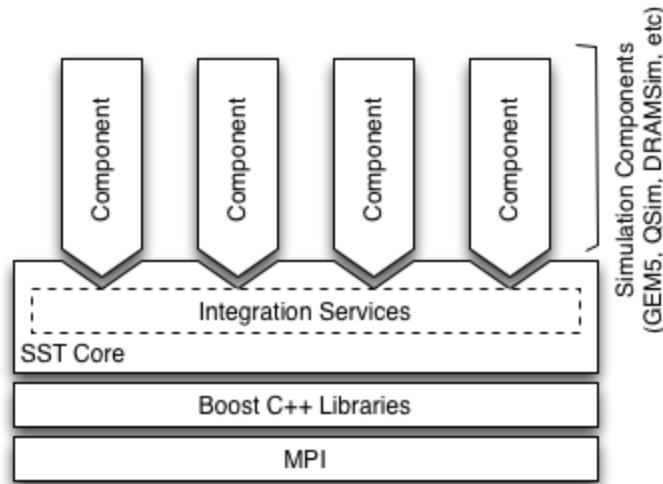
- Propose
 - Exploring novel computer system designs
- problems
 - serial
 - lack of scalability and flexibility
- Technical Approach
 - compiler tools, well defined ISA features and complete microarchitectural definitions ...
 - Parallel - MPI + Threads
 - Flexible - simulation components
 - Open API - easily extensible with new models

Existing SST Element Libraries



Detailed Memory Models	memHierarchy - Cache and Memory
	cassini - Cache prefetchers
	DRAMSim2 - DDR
	NVDIMMSim - Emerging Memories
	Goblin - HMC
	ariel - PIN-based Tracing
Dynamic Trace-based Processor Model	MacSim - GPGPU
Cycle-based Processor Model	m5C - Gem5 integration layer
High-level Program Communication Models	ember - State-machine Message generation
	firefly - Communication Protocols
	hermes - MPI-like interface
Cycle-based Network Model	merlin - Network router model and NIC
High-level System Workflow Model	scheduler - Job-scheduler simulation models

- component

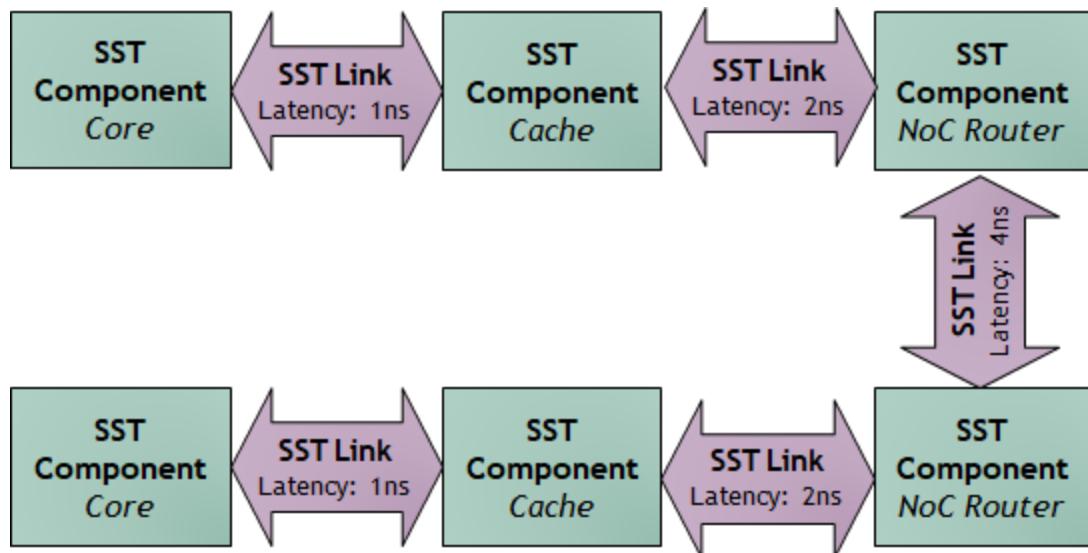


- core - contain the simulation engine and API interface to simulation elements
 - clocks, event exchange, statistics, parameter management, parallelism support

- elements
 - memory, cache, processor architecture, networking ...
- macro - functionality for simulating extreme-scale applications

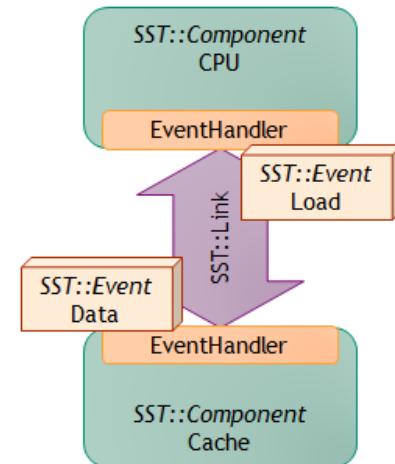
High-level View

- comprised of components connected by ***links***
- Components define ***ports*** which are valid connection points for a link
- Components can use ***subComponents*** and ***modules*** for extra functionality



12 Key objects

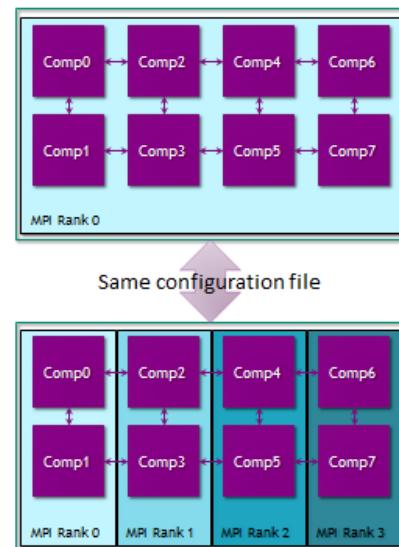
- **SST::Component**
 - Simulation model
- **SST::Link**
 - Communication path between two components
 - Poll or interrupt with EventHandler
- **SST::Event**
 - A discrete event
- **SST::Clock::Handler**
 - Function to handle a clock tick
- **SST::SubComponent**
 - Add functionality to Components



- **Polled:** Register a clock handler to poll the link
- **Interrupt:** Register an event handler to be called when an event arrives

22 SST in Parallel

- SST was designed from the ground up to enable scalable, parallel simulations
- Components distributed among MPI ranks
 - And/or threads
- Links allow parallelism
 - Hence, components should communicate via links only
 - Transparently handle any MPI/thread communication
 - Specified link-latency determines synchronization rate



Case Study

19

Simulating with SST

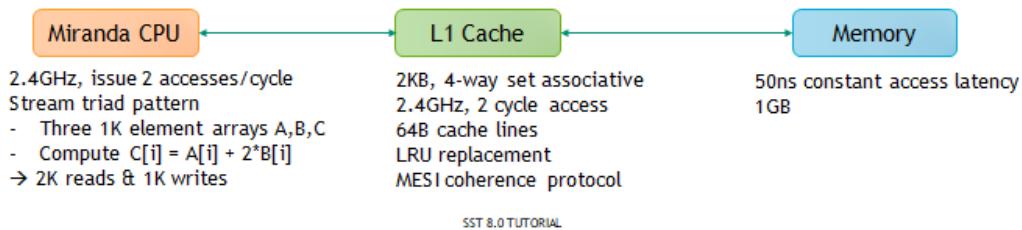


- To define simulations, users create a python configuration file which:
 - Defines the components and subcomponents to use
 - Specifies the connectivity between components
- SST then:
 - Loads the file and converts it to a graph
 - Partitions the graph for parallel simulation
 - Instantiates the components
 - Runs the simulation!

SST 8.0 TUTORIAL

26 | Running a simulation: Example

- We'll walk through the configuration file, then run it
- Full intro to the elements is next but until then this example uses:
 - **Miranda** – Pattern-based memory access generator
 - With a *stream* memory access pattern generator
 - **MemHierarchy** – Caches and memory models
 - Just an L1 cache and memory for this example



27 | Configuration Part I: Configure SST

- Set any global simulation parameters
- `sst.setProgramOption("stopAtCycle", "100ms")`
 - Kill simulation (nicely!) if it runs to 100ms

SST 8.0 TUTORIAL

28 Configuration Part 2: Declare components

◦ **Define:** `sst.Component("name", "type")`

```
core = sst.Component("core", "miranda.BaseCPU")
cache = sst.Component("L1", "memHierarchy.Cache")
memory = sst.Component("memory", "memHierarchy.MemController")
```

Component name

Component type

◦ **Configure:** `addParams({ "parameter" : "value", ... })`

```
core.addParams({
    "clock" : "2.4GHz",
    "generator" : "miranda.STREAMBenchGenerator",
    "generator.n" : 1000,
})
...
# cache and memory parameters
```

Component parameters

SST 8.0 TUTORIAL

30 Configuration Part 3: Connect the components

◦ **Create a link:** `sst.Link("name")`

```
core_cache = sst.Link("core_to_cache")
cache_mem = sst.Link("cache_to_memory")
```

Link name
(anything unique)

◦ **Connect components:** `connect(endpoint1, endpoint2)`

◦ Where endpoint is: (component, port, latency)

```
core_cache.connect(
    (core, "cache_link", "100ps"),
    (cache, "high_network_0", "100ps" ) )
cache_mem.connect(
    (cache, "low_network_0", "100ps"),
    (memory, "direct_link", "100ps" ) )
```

Endpoint 1

Endpoint 2

SST 8.0 TUTORIAL

- Language

- Python

- C/C++
- Required Tools(?)
 - MPI
 - Boost C++ Libraries
- Critics?
 - <https://cfwebprod.sandia.gov/cfdocs/CompResearch/docs/index1.pdf>

Table 1. Simulation Parameters

Element	Parameters
CPU	8 Out-of-Order cores, 2GHz, 2 issues/cycles, 32 max. outstanding requests
L1	private, 64B blocks, 32KB, LRU
L2	private, 64B blocks, 256KB, LRU
L3	shared, 64B blocks ,16MB, LRU
Local memory	2GB, DDR4-based DRAM
Global memory	16GB, NVM-based DIMM (PCM), 128 max. outstanding requests, 16 banks, 300ns Read Latency, 1000ns Write Latency
External network latency	20ns[4]

- Potential Direction
 - developing SST Element Libraries
 -

References

References



- **Websites**
 - <http://www.sst-simulator.org/>
 - <https://github.com/sstsimulator>
- **Configuration File Format:**
 - <http://sst-simulator.org/SSTPages/SSTUserPythonFileFormat/>
- **Doxygen Documentation:**
 - http://sst-simulator.org/SSTDoxygen/8.0.0_docs/html/
- **Developer FAQ:**
 - <http://sst-simulator.org/SSTPages/SSTTopDocDeveloperInfo/>
- **Building SST**
 - <http://sst-simulator.org/SSTPages/SSTBuildAndInstall8dot0dot0SeriesQuickStart/>
 - <http://sst-simulator.org/SSTPages/SSTBuildAndInstall8dot0dot0SeriesDetailedBuildInstructions/>

SST 8.0 TUTORIAL

https://hpc.pnl.gov//modsim/2017/Presentations/SST_Modsim17.pdf

SST Publications

<http://sst-simulator.org/SSTPages/SSTTopDocPublications/>